

k -Means Clustering

Julia Kempe & David S. Rosenberg

CDS, NYU


May 1, 2019

Contents

- 1 k -Means Clustering
- 2 k -Means: Failure Cases
- 3 k -means Formalized

k -Means Clustering


Example: Old Faithful Geyser



- Looks like two clusters.
- How to find these clusters algorithmically?

k-Means: By Example


- Standardize the data.
- Choose two cluster centers.



From Bishop's *Pattern recognition and machine learning*, Figure 9.1(a).

k-means: by example


- Assign each point to closest center.



From Bishop's *Pattern recognition and machine learning*, Figure 9.1(b).

k -means: by example


- Compute new class centers.



From Bishop's *Pattern recognition and machine learning*, Figure 9.1(c).

k-means: by example


- Assign points to closest center.



From Bishop's *Pattern recognition and machine learning*, Figure 9.1(d).

k -means: by example


- Compute cluster centers.



From Bishop's *Pattern recognition and machine learning*, Figure 9.1(e).

k -means: by example


- Iterate until convergence.



From Bishop's *Pattern recognition and machine learning*, Figure 9.1(i).

k-Means Algorithm: Standardizing the data


- Without standardizing:



- Blue and black show results of k-means clustering
- Wait time dominates the distance metric

k-Means Algorithm: Standardizing the data

- With standardizing:



- Note several points have been reassigned from black to blue cluster


k -Means: Failure Cases

k -Means: Suboptimal Local Minimum

- The clustering for $k = 3$ below is a local minimum, but suboptimal:



Would be better to have
one cluster here



... and two clusters here

k -means Formalized

k -Means: Setting

- Let \mathcal{X} be a space with some distance metric d .
 - Most commonly, $\mathcal{X} = \mathbf{R}^d$ and $d(x, x') = \|x - x'\|$.
- Dataset $\mathcal{D} = \{x_1, \dots, x_n\} \subset \mathcal{X}$.
- Goal: Partition data \mathcal{D} into k disjoint sets C_1, \dots, C_k .
- The **centroid** of C_i is defined to be

$$\mu_i = \mu(C_i) = \arg \min_{\mu \in \mathcal{X}} \sum_{x \in C_i} d(x, \mu)^2.$$

- Note: For Euclidean distance on \mathbf{R}^d , $\mu(C_i)$ is the mean of C_i .

k -Means: Objective function

- The k -means objective is

$$\begin{aligned} J_{k\text{-means}}(C_1, \dots, C_k) &= \sum_{i=1}^k \sum_{x \in C_i} d(x, \mu(C_i))^2 \\ &= \min_{\mu_1, \dots, \mu_k \in \mathcal{X}} \sum_{i=1}^k \sum_{x \in C_i} d(x, \mu_i)^2 \end{aligned}$$

- In **vector quantization**, we represent each $x \in C_i$ by the centroid μ_i .
- We can think of this as lossy data compression,
 - the k -means objective can be viewed as the reconstruction error.
- How many bits does it take to represent each point with vector quantization?
 - If $k = 2^d$, then d bits. (Fewer on average if the clusters have unequal sizes.)

k -Means: Algorithm

- input: $\mathcal{D} = \{x_1, \dots, x_d\} \subset \mathcal{X}$
- initialize: Randomly choose initial centroids $\mu_1, \dots, \mu_k \subset \mathcal{D}$.
- repeat until convergence (i.e. until the centroids or clusters repeat):
 - $\forall i$, let $C_i = \{x \in \mathcal{D} : i = \arg \min_j d(x, \mu_j)\}$. (break ties in some arbitrary manner)
 - $\forall i$, let $\mu_i = \arg \min_{\mu \in \mathcal{X}} \sum_{x \in C_i} d(x, \mu)^2$. (For Euclidean distance, $\mu_i = \frac{1}{|C_i|} \sum_{x \in C_i} x$)

- In k -means, objective never increases, but no guarantee to find minimizer.
- General recommendation is to re-run with several random starting initial centroids.
- k -means++ is a way to randomly initialize the centroids with some guarantees:
 - Randomly choose first centroid from the data points \mathcal{D} .
 - For each of the remaining $k-1$ centroids:
 - Compute distance from each x_i to the closest already chosen centroid.
 - Randomly choose next centroid with probability proportional to the computed distance squared.
- If we let $J_{k\text{-means}}^*$ be the minimizer of the k -means objective, then using k -means++ for initialization guarantees that

$$\mathbb{E}[J_{k\text{-means}}(C_1, \dots, C_k)] \leq 8(\log k + 2) J_{k\text{-means}}^*.$$