

# Gaussian Mixture Models

Julia Kempe & David S. Rosenberg

CDS, NYU

May 7, 2019

# Contents

- 1 Gaussian Mixture Models
- 2 Mixture Models
- 3 Learning in Gaussian Mixture Models
- 4 Issues with MLE for GMM
- 5 The EM Algorithm for GMM


# Gaussian Mixture Models

# Probabilistic Model for Clustering

- Let's consider the following **generative model** (i.e. a way to generate data).
- Suppose
  - ① There are  $k$  clusters (or “**mixture components**”).
  - ② We have a probability density for each cluster.
- Generate a point as follows
  - ① Choose a random cluster  $z \in \{1, 2, \dots, k\}$ .
  - ② Choose a point from the distribution for cluster  $z$ .
- Data generated in this way is said to have a **mixture distribution**.

# Gaussian Mixture Model ( $k = 3$ )

- 1 Choose  $z \in \{1, 2, 3\}$  with  $p(1) = p(2) = p(3) = \frac{1}{3}$ .
- 2 Choose  $x \mid z \sim \mathcal{N}(X \mid \mu_z, \Sigma_z)$ .




# Gaussian Mixture Model Parameters ( $k$ Components)

Cluster probabilities:  $\pi = (\pi_1, \dots, \pi_k)$

Cluster means:  $\mu = (\mu_1, \dots, \mu_k)$

Cluster covariance matrices:  $\Sigma = (\Sigma_1, \dots, \Sigma_k)$



For now, **suppose all these parameters are known.**  
We'll discuss how to **learn** or **estimate** them later.

# Gaussian Mixture Model: Joint Distribution

- Factorize the joint density:

$$\begin{aligned} p(x, z) &= p(z)p(x | z) \\ &= \pi_z \mathcal{N}(x | \mu_z, \Sigma_z) \end{aligned}$$

- $\pi_z$  is probability of choosing cluster  $z$ .
- $x | z$  has distribution  $\mathcal{N}(\mu_z, \Sigma_z)$ .
- $z$  corresponding to  $x$  is the true cluster assignment.
- Suppose we know the model parameters  $\pi_z, \mu_z, \Sigma_z$ .
- Then we can easily evaluate the joint density  $p(x, z)$ .

# Latent Variable Model

- We observe  $x$ .
- We don't observe  $z$  (the cluster assignment).
- Cluster assignment  $z$  is called a **hidden variable** or **latent variable**.

## Definition

A **latent variable model** is a probability model for which certain variables are never observed.

e.g. The Gaussian mixture model is a latent variable model.



# The GMM “Inference” Problem

- We observe  $x$ . We want to know its cluster assignment  $z$ .
- The conditional probability for cluster  $z$  given  $x$  is

$$p(z | x) = p(x, z) / p(x)$$

- The conditional distribution is a **soft assignment** to clusters.
- A **hard assignment** is

$$z^* = \arg \max_{z \in \{1, \dots, k\}} p(z | x).$$

- So if we know the model parameters, clustering is trivial.

# Mixture Models

# General Mixture Models: Generative Construction

- Let  $S$  be a set of  $k$  probability distributions (“**mixture components**”).
- Let  $\pi = (\pi_1, \dots, \pi_k)$  be a distribution on  $\{1, \dots, k\}$  (“**mixture weights**”).
- Suppose we generate  $x$  with the following procedure:
  - ① Choose a distribution randomly from  $S$  according to  $\pi$ .
  - ② Sample  $x$  from the chosen distribution.
- Then we say  $x$  has a **mixture distribution**.

# Mixture Densities

- Suppose we have a mixture distribution with
  - **mixture components** represented as densities  $p_1, \dots, p_k$ , and
  - **mixture weights**  $\pi = (\pi_1, \dots, \pi_k)$ , then
- the corresponding probability density for  $x$  is

$$p(x) = \sum_{i=1}^k \pi_i p_i(x).$$

- Note that  $p$  is a convex combination of the mixture component densities.
- $p(x)$  is called a **mixture density**.
- Conversely, if  $x$  has a density of this form, then  $x$  has a mixture distribution.

# Gaussian Mixture Model (GMM): Marginal Distribution

For example:

- The **marginal distribution** for a single observation  $x$  in a GMM is

$$\begin{aligned} p(x) &= \sum_{z=1}^k p(x, z) \\ &= \sum_{z=1}^k \pi_z \mathcal{N}(x \mid \mu_z, \Sigma_z) \end{aligned}$$

# Learning in Gaussian Mixture Models

---

# The GMM “Learning” Problem

- Given data  $x_1, \dots, x_n$  drawn from a GMM,
- Estimate the parameters:

Cluster probabilities:  $\pi = (\pi_1, \dots, \pi_k)$

Cluster means:  $\mu = (\mu_1, \dots, \mu_k)$

Cluster covariance matrices:  $\Sigma = (\Sigma_1, \dots, \Sigma_k)$

- Once we have the parameters, we're done.
- Just do “inference” to get cluster assignments.

# Estimating/Learning the Gaussian Mixture Model

- One approach to learning is **maximum likelihood**
  - find parameter values with **highest likelihood** for the **observed data**.
- The model likelihood for  $\mathcal{D} = (x_1, \dots, x_n)$  sampled iid from a GMM is

$$\begin{aligned} L(\pi, \mu, \Sigma) &= \prod_{i=1}^n p(x_i) \\ &= \prod_{i=1}^n \sum_{z=1}^k \pi_z \mathcal{N}(x_i \mid \mu_z, \Sigma_z). \end{aligned}$$

- As usual, we'll take our objective function to be the log of this:

$$J(\pi, \mu, \Sigma) = \sum_{i=1}^n \log \left\{ \sum_{z=1}^k \pi_z \mathcal{N}(x_i \mid \mu_z, \Sigma_z) \right\}$$



## Review: Estimating a Gaussian Distribution

- Recall that the density for  $x \sim \mathcal{N}(\mu, \Sigma)$  is

$$p(x \mid \mu, \Sigma) = \frac{1}{\sqrt{|2\pi\Sigma|}} \exp\left(-\frac{1}{2}(x - \mu)^T \Sigma^{-1}(x - \mu)\right)$$

- And the log-density is

$$\log p(x \mid \mu, \Sigma) = -\frac{1}{2} \log |2\pi\Sigma| - \frac{1}{2}(x - \mu)^T \Sigma^{-1}(x - \mu)$$

- To estimate  $\mu$  and  $\Sigma$  from a sample  $x_1, \dots, x_n$  i.i.d.  $\mathcal{N}(\mu, \Sigma)$ , we'll maximize the log joint density:

$$\sum_{i=1}^n \log p(x_i \mid \mu, \Sigma) = -\frac{n}{2} \log |2\pi\Sigma| - \frac{1}{2} \sum_{i=1}^n (x_i - \mu)^T \Sigma^{-1}(x_i - \mu)$$

## Review: Estimating a Gaussian Distribution

- To estimate  $\mu$  and  $\Sigma$  from a sample  $x_1, \dots, x_n$  i.i.d.  $\mathcal{N}(\mu, \Sigma)$ , we'll maximize the log joint density:

$$J(\mu, \Sigma) = \sum_{i=1}^n \log p(x_i | \mu, \Sigma) = -\frac{n}{2} \log |2\pi\Sigma| - \frac{1}{2} \sum_{i=1}^n (x_i - \mu)^T \Sigma^{-1} (x_i - \mu)$$

- This is a solid exercise in vector and matrix differentiation. Find  $\hat{\mu}$  and  $\hat{\Sigma}$  satisfying

$$\nabla_{\mu} J(\mu, \Sigma) = 0 \quad \nabla_{\Sigma} J(\mu, \Sigma) = 0$$

- We get a closed form solution:

$$\begin{aligned} \hat{\mu}_{\text{MLE}} &= \frac{1}{n} \sum_{i=1}^n x_i \\ \hat{\Sigma}_{\text{MLE}} &= \frac{1}{n} \sum_{i=1}^n (x_i - \hat{\mu}_{\text{MLE}})^T (x_i - \hat{\mu}_{\text{MLE}}) \end{aligned}$$

# Properties of the GMM Log-Likelihood

- GMM log-likelihood:

$$J(\pi, \mu, \Sigma) = \sum_{i=1}^n \log \left\{ \sum_{z=1}^k \frac{\pi_z}{\sqrt{|2\pi\Sigma_z|}} \exp \left( -\frac{1}{2} (x - \mu_z)^T \Sigma_z^{-1} (x - \mu_z) \right) \right\}$$

- Let's compare to the log-likelihood for a single Gaussian:

$$J(\mu, \Sigma) = -\frac{n}{2} \log |2\pi\Sigma| - \frac{1}{2} \sum_{i=1}^n (x_i - \mu)^T \Sigma^{-1} (x_i - \mu)$$

- For a single Gaussian, the log cancels the exp in the Gaussian density.
  - $\implies$  Things simplify a lot.
- For the GMM, the sum inside the log prevents this cancellation.
  - $\implies$  Expression more complicated. No closed form expression for MLE.

## Issues with MLE for GMM

# Identifiability Issues for GMM

- Suppose we have found parameters

Cluster probabilities:  $\pi = (\pi_1, \dots, \pi_k)$

Cluster means:  $\mu = (\mu_1, \dots, \mu_k)$


Cluster covariance matrices:  $\Sigma = (\Sigma_1, \dots, \Sigma_k)$

that are at a local minimum.

- What happens if we shuffle the clusters? e.g. Switch the labels for clusters 1 and 2.
- We'll get the same likelihood. How many such equivalent settings are there?
- Assuming all clusters are distinct, there are  $k!$  equivalent solutions.
- Not a problem *per se*, but something to be aware of.

# Singularities for GMM

- Consider the following GMM for 7 data points:



- Let  $\sigma^2$  be the variance of the skinny component.
- What happens to the likelihood as  $\sigma^2 \rightarrow 0$ ?
- In practice, we end up in local minima that do not have this problem.
  - Or keep restarting optimization until we do.
- Bayesian approach or regularization will also solve the problem.

From Bishop's *Pattern recognition and machine learning*, Figure 9.7.

# Gradient Descent / SGD for GMM

- What about running gradient descent or SGD on

$$J(\pi, \mu, \Sigma) = - \sum_{i=1}^n \log \left\{ \sum_{z=1}^k \pi_z \mathcal{N}(x_i \mid \mu_z, \Sigma_z) \right\}?$$

- Can be done, in principle – but need to be clever about it.
- Each matrix  $\Sigma_1, \dots, \Sigma_k$  has to be positive semidefinite.
- How to maintain that constraint?
  - Rewrite  $\Sigma_i = M_i M_i^T$ , where  $M_i$  is an unconstrained matrix.
  - Then  $\Sigma_i$  is positive semidefinite.
- Even then, pure gradient-based methods have trouble.<sup>1</sup>

---

<sup>1</sup>See Hosseini and Sra's [Manifold Optimization for Gaussian Mixture Models](#) for discussion and further references.

# The EM Algorithm for GMM

---



# MLE for Gaussian Model

- Let's start by considering the MLE for the Gaussian model.
- For data  $\mathcal{D} = \{x_1, \dots, x_n\}$ , the log likelihood is given by

$$\sum_{i=1}^n \log \mathcal{N}(x_i | \mu, \Sigma) = -\frac{nd}{2} \log(2\pi) - \frac{n}{2} \log |\Sigma| - \frac{1}{2} \sum_{i=1}^n (x_i - \mu)' \Sigma^{-1} (x_i - \mu).$$

- With some calculus, we find that the MLE parameters are

$$\begin{aligned}\mu_{\text{MLE}} &= \frac{1}{n} \sum_{i=1}^n x_i \\ \Sigma_{\text{MLE}} &= \frac{1}{n} \sum_{i=1}^n (x_i - \mu_{\text{MLE}})(x_i - \mu_{\text{MLE}})^T\end{aligned}$$

- For GMM, If we knew the cluster assignment  $z_i$  for each  $x_i$ ,
  - we could compute the MLEs for each cluster.

# Estimating a Fully-Observed GMM

- Suppose we observe  $(x_1, z_1), \dots, (x_n, z_n)$  i.i.d. from GMM  $p(x, z)$ .
- Then find MLE is easy:

$$\begin{aligned}n_z &= \sum_{i=1}^n 1(z_i = z) \\ \hat{\pi}(z) &= \frac{n_z}{n} \\ \hat{\mu}_z &= \frac{1}{n_z} \sum_{i: z_i = z} x_i \\ \hat{\Sigma}_z &= \frac{1}{n_z} \sum_{i: z_i = z} (x_i - \hat{\mu}_z)(x_i - \hat{\mu}_z)^T.\end{aligned}$$

- In the EM algorithm we will modify the equations to handle our evolving **soft assignments**, which we will call **responsibilities**.

## Cluster Responsibilities: Some New Notation

- Denote the probability that observed value  $x_i$  comes from cluster  $j$  by

$$\gamma_i^j = p(z = j \mid x = x_i).$$

- The **responsibility** that cluster  $j$  takes for observation  $x_i$ .
- Computationally,

$$\begin{aligned}\gamma_i^j &= p(z = j \mid x_i) \\ &= p(z = j, x_i) / p(x_i) \\ &= \frac{\pi_j \mathcal{N}(x_i \mid \mu_j, \Sigma_j)}{\sum_{c=1}^k \pi_c \mathcal{N}(x_i \mid \mu_c, \Sigma_c)}\end{aligned}$$

- The vector  $(\gamma_i^1, \dots, \gamma_i^k)$  is exactly the **soft assignment** for  $x_i$ .
- Let  $n_c = \sum_{i=1}^n \gamma_i^c$  be the “number” of points “soft assigned” to cluster  $c$ .

# EM Algorithm for GMM: Overview

- If we know  $\mu_j, \Sigma_j, \pi_j$  for all clusters  $j$ , then easy to find

$$\gamma_i^j = p(z=j \mid x_i)$$

- If we know the (soft) assignments, we can easily find estimates for  $\pi, \Sigma, \mu$ .
- Repeatedly alternate these two steps.

# EM Algorithm for GMM: Overview

- 1 Initialize parameters  $\mu, \Sigma, \pi$  (e.g. using  $k$ -means).
- 2 “E step”. Evaluate the responsibilities using current parameters:

$$\gamma_i^j = \frac{\pi_j \mathcal{N}(x_i | \mu_j, \Sigma_j)}{\sum_{c=1}^k \pi_c \mathcal{N}(x_i | \mu_c, \Sigma_c)},$$

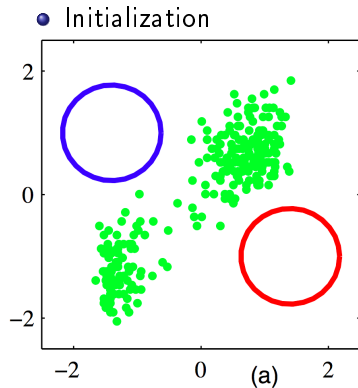
for  $i = 1, \dots, n$  and  $j = 1, \dots, k$ .

- 3 “M step”. Re-estimate the parameters using responsibilities:

$$\begin{aligned}\mu_c^{\text{new}} &= \frac{1}{n_c} \sum_{i=1}^n \gamma_i^c x_i \\ \Sigma_c^{\text{new}} &= \frac{1}{n_c} \sum_{i=1}^n \gamma_i^c (x_i - \mu_c^{\text{new}}) (x_i - \mu_c^{\text{new}})^T \\ \pi_c^{\text{new}} &= \frac{n_c}{n},\end{aligned}$$

- 4 Repeat from Step 2, until log-likelihood converges.


# EM for GMM



From Bishop's *Pattern recognition and machine learning*, Figure 9.8.

# EM for GMM


• First soft assignment:



From Bishop's *Pattern recognition and machine learning*, Figure 9.8.

# EM for GMM

- First soft assignment:




From Bishop's *Pattern recognition and machine learning*, Figure 9.8.



# EM for GMM


• After 5 rounds of EM:



From Bishop's *Pattern recognition and machine learning*, Figure 9.8.

# EM for GMM

• After 20 rounds of EM:



From Bishop's *Pattern recognition and machine learning*, Figure 9.8.

## Relation to $k$ -Means

- EM for GMM seems a little like  $k$ -means.
- In fact,  $k$ -means is a limiting case of a **restricted** version of GMM.
- First, fix each cluster covariance matrix to be  $\sigma^2 I$ .
  - (This is the restriction: covariance matrices are fixed, and not iteratively estimated.)
- As we take  $\sigma^2 \rightarrow 0$ , the update equations converge to doing  $k$ -means.
- If you do a quick experiment yourself, you'll find
  - Soft assignments converge to hard assignments.
  - Has to do with the tail behavior (exponential decay) of Gaussian.